

SQL DW Management and Monitoring (ADW In-A-Day Lab 03)

Contents

SQL DW Management and Monitoring (ADW In-A-Day Lab 03)	1
Overview	1
Pre-requisites	1
Resource Monitoring in Azure Monitor	1
Azure Data Studio SQLDW Dashboard (Azure SQL Data Warehouse Insights).....	2
Azure SQL Data Warehouse Table and Statistics Queries	3
Create User-defined Restore Points	7
Maintenance Window Scheduling, Service Health, Service Health Alerts	14
Querying ADW Diagnostic Logs using Azure Monitor.....	18

Overview


This module will walk you through a variety of tools and techniques for monitoring and managing your Azure Data Warehouse.

Pre-requisites

- Azure Portal Access
- Azure SQL Data Warehouse
- SQL Server Management Studio
- Azure Data Studio

Resource Monitoring in Azure Monitor

1. Logon to Azure Portal (portal.azure.com) using your credentials
2. Click on Monitor from the list of services.





3. Click on 'Explore Metrics'
4. Click on 'Add metric' and select Subscription and Resource Group names from the drop-down.
5. Select your DW instance from the list of resources displayed.
6. Select 'CPU percentage' as the metric from the drop-down as shown below:

RESOURCE	METRIC NAMESPACE	METRIC	AGGREGATION
 kalsql/kaldw2	Sql database standa...	CPU percentage	Avg

- This will display the CPU percentage (Avg) chart for the last 24 hours.
- Click on 'Add metric' again and 'Data IO Percentage' to the chart.
- Repeat the process to add more metrics to the chart.
- Click on the Time Range drop-down and change the time range from 'Last 24 hours' to 'Last hour'. Observe the chart data range change to display last hour data.
- Click on 'Add chart' and select Subscription, Resource Group, Resource and Metric as before. You will observe a new chart created displaying the data for the metric you selected.

Azure Data Studio SQLDW Dashboard (Azure SQL Data Warehouse Insights)

- Open Azure Data Studio application.
Azure Data Studio (ADS) is a cross-platform database tool for monitoring and managing various data platforms, including Azure SQL DW.
- Click on 'Extensions'  icon from the list of icons on the left side ribbon of the window.
- Select 'Azure SQL Data Warehouse Insights' from the list of extensions. (Note: Enter 'SQL' in the search window, if this extension is not showing up in the list)
- Click 'Install' and wait for the extension installation to complete. (Note: Review the details of the extension to understand all the reports displayed by this extension. You can also click on the github link available to review the sql code for each of these reports)
- Click on the 'Servers'  icon on the left and click on 'Add Connection'.
- Expand 'Databases' node in the explorer window and right-click on your database and click 'Manage'. This will show Database Dashboard with details on your DW instance and tables.
- Click on 'SQL DW Dashboard' tab in the details pane. Observe the reports on various metrics viz. Latest Backup date, User Activities, Data Distribution etc. (Note: Some reports maybe empty if there are no results to display.)
- Observe the Data Distribution report. This shows data size in each of the 60 distributions. What can you infer from this report?

Azure SQL Data Warehouse Table and Statistics Queries

1. Open SQL Server Management Studio (SSMS). Connect to the your SQL DW instance.
2. Expand Databases node, right-click on your data warehouse database and click 'New Query'. Execute the following query in your query window to create a useful View.

```

CREATE VIEW dbo.vTableSizes
AS
WITH base
AS
(
SELECT
    GETDATE() AS [execution_time]
    , DB_NAME() AS [database_name]
    , s.name AS [schema_name]
    , t.name AS [table_name]
    , QUOTENAME(s.name)+'.'+QUOTENAME(t.name) AS [two_part_name]
    , nt.name AS [node_table_name]
    , ROW_NUMBER() OVER(PARTITION BY nt.[name] ORDER BY (SELECT NULL)) AS
[node_table_name_seq]
    , tp.[distribution_policy_desc] AS
[distribution_policy_name]
    , c.[name] AS [distribution_column]
    , nt.[distribution_id] AS [distribution_id]
    , i.[type] AS [index_type]
    , i.[type_desc] AS [index_type_desc]
    , nt.[pdw_node_id] AS [pdw_node_id]
    , pn.[type] AS [pdw_node_type]
    , pn.[name] AS [pdw_node_name]
    , di.name AS [dist_name]
    , di.position AS [dist_position]
    , nps.[partition_number] AS [partition_nmbr]
    , nps.[reserved_page_count]
    , nps.[reserved_page_count] - nps.[used_page_count] AS
[unused_space_page_count]
    , nps.[in_row_data_page_count]
    + nps.[row_overflow_used_page_count]
    + nps.[lob_used_page_count] AS [data_space_page_count]
    , nps.[reserved_page_count]
    - (nps.[reserved_page_count] - nps.[used_page_count])
    - ([in_row_data_page_count]
    + [row_overflow_used_page_count]+[lob_used_page_count]) AS
[index_space_page_count]
    , nps.[row_count] AS [row_count]
from
    sys.schemas s
INNER JOIN sys.tables t
    ON s.[schema_id] = t.[schema_id]
INNER JOIN sys.indexes i
    ON t.[object_id] = i.[object_id]
    AND i.[index_id] <= 1
INNER JOIN sys.pdw_table_distribution_properties tp
    ON t.[object_id] = tp.[object_id]
INNER JOIN sys.pdw_table_mappings tm
    ON t.[object_id] = tm.[object_id]
INNER JOIN sys.pdw_nodes_tables nt
    ON tm.[physical_name] = nt.[name]
INNER JOIN sys.dm_pdw_nodes pn
    ON nt.[pdw_node_id] = pn.[pdw_node_id]
INNER JOIN sys.pdw_distributions di
    ON nt.[distribution_id] = di.[distribution_id]

```

```

INNER JOIN sys.dm_pdw_nodes_db_partition_stats nps
  ON nt.[object_id] = nps.[object_id]
  AND nt.[pdw_node_id] = nps.[pdw_node_id]
  AND nt.[distribution_id] = nps.[distribution_id]
LEFT OUTER JOIN (select * from sys.pdw_column_distribution_properties where distribution_ordinal =
1) cdp
  ON t.[object_id] = cdp.[object_id]
LEFT OUTER JOIN sys.columns c
  ON cdp.[object_id] = c.[object_id]
  AND cdp.[column_id] = c.[column_id]
)
, size
AS
(
SELECT
  [execution_time]
, [database_name]
, [schema_name]
, [table_name]
, [two_part_name]
, [node_table_name]
, [node_table_name_seq]
, [distribution_policy_name]
, [distribution_column]
, [distribution_id]
, [index_type]
, [index_type_desc]
, [pdw_node_id]
, [pdw_node_type]
, [pdw_node_name]
, [dist_name]
, [dist_position]
, [partition_nmbr]
, [reserved_space_page_count]
, [unused_space_page_count]
, [data_space_page_count]
, [index_space_page_count]
, [row_count]
, ([reserved_space_page_count] * 8.0) AS
[reserved_space_KB]
, ([reserved_space_page_count] * 8.0)/1000 AS
[reserved_space_MB]
, ([reserved_space_page_count] * 8.0)/1000000 AS
[reserved_space_GB]
, ([reserved_space_page_count] * 8.0)/1000000000 AS
[reserved_space_TB]
, ([unused_space_page_count] * 8.0) AS
[unused_space_KB]
, ([unused_space_page_count] * 8.0)/1000 AS
[unused_space_MB]
, ([unused_space_page_count] * 8.0)/1000000 AS
[unused_space_GB]
, ([unused_space_page_count] * 8.0)/1000000000 AS
[unused_space_TB]
, ([data_space_page_count] * 8.0) AS
[data_space_KB]
, ([data_space_page_count] * 8.0)/1000 AS
[data_space_MB]
, ([data_space_page_count] * 8.0)/1000000 AS
[data_space_GB]
, ([data_space_page_count] * 8.0)/1000000000 AS
[data_space_TB]

```

```

, ([index_space_page_count] * 8.0) AS
[index_space_KB]
, ([index_space_page_count] * 8.0)/1000 AS
[index_space_MB]
, ([index_space_page_count] * 8.0)/1000000 AS
[index_space_GB]
, ([index_space_page_count] * 8.0)/1000000000 AS
[index_space_TB]
FROM base
)
SELECT *
FROM size
;
GO

```

9. Execute the following query in the query window. This will show table space summary information.

```

SELECT
    database_name
,    schema_name
,    table_name
,    distribution_policy_name
,    distribution_column
,    index_type_desc
,    COUNT(distinct partition_nmbr) as nbr_partitions
,    SUM(row_count) as table_row_count
,    SUM(reserved_space_GB) as table_reserved_space_GB
,    SUM(data_space_GB) as table_data_space_GB
,    SUM(index_space_GB) as table_index_space_GB
,    SUM(unused_space_GB) as table_unused_space_GB
FROM
    dbo.vTableSizes
GROUP BY
    database_name
,    schema_name
,    table_name
,    distribution_policy_name
,    distribution_column
,    index_type_desc
ORDER BY
    table_reserved_space_GB desc

```

10. Click 'New Query'. Execute the following query in your query window. This will list tables with >10% skew.

```

select two_part_name
,[distribution_column]
,(max(row_count * 1.000) - min(row_count * 1.000))/max(row_count * 1.000) as
'skew'
,max(row_count) as 'largest dist'
,min(row_count) as 'smallest dist'
,avg(row_count) as 'average dist'
,sum(row_count) as 'total row_count'
from dbo.vTableSizes
where row_count > 0
AND distribution_policy_name = 'HASH'
group by two_part_name, [distribution_column]
having (max(row_count * 1.000) - min(row_count * 1.000))/max(row_count * 1.000)
>= .10
order by 3 DESC -- two_part_name

```

3. Expand Databases node, right-click on your data warehouse database and click 'New Query'. Execute the following query in your query window to create another useful View.

```

CREATE VIEW dbo.vStats_Columns
AS
SELECT
    sm.[name] AS [schema_name]
    , tb.[name] AS [table_name]
    , st.[name] AS [stats_name]
    , st.[filter_definition] AS [stats_filter_definition]
    , st.[has_filter] AS [stats_is_filtered]
    , STATS_DATE(st.[object_id], st.[stats_id]) AS [stats_last_updated_date]
    , st.[user_created] AS [user_created]
    , co.[name] AS [stats_column_name]
    , ty.[name] AS [column_type]
    , co.[max_length] AS [column_max_length]
    , co.[precision] AS [column_precision]
    , co.[scale] AS [column_scale]
    , co.[is_nullable] AS [column_is_nullable]
    , co.[collation_name] AS [column_collation_name]
    , QUOTENAME(sm.[name])+'.'+QUOTENAME(tb.[name]) AS two_part_name
    , QUOTENAME(DB_NAME())+'.'+QUOTENAME(sm.[name])+'.'+QUOTENAME(tb.[name]) AS three_part_name
    , QUOTENAME(sm.[name])+'.'+QUOTENAME(tb.[name])+'.'+QUOTENAME(st.[name]) AS full_stats_name
    , st.[stats_generation_method_desc] AS [stats_generation_method_desc]
FROM sys.objects AS ob
JOIN sys.stats AS st ON ob.[object_id] = st.[object_id]
JOIN sys.stats_columns AS sc ON st.[stats_id] = sc.[stats_id]
    AND st.[object_id] = sc.[object_id]
JOIN sys.columns AS co ON sc.[column_id] = co.[column_id]
    AND sc.[object_id] = co.[object_id]
JOIN sys.types AS ty ON co.[user_type_id] = ty.[user_type_id]
JOIN sys.tables AS tb ON co.[object_id] = tb.[object_id]
JOIN sys.schemas AS sm ON tb.[schema_id] = sm.[schema_id]
;
GO

```

11. Execute the following query in the query window. This will show table statistics information.

```
SELECT * FROM dbo.vStats_Columns ORDER BY 1,2,3
```


Create User-defined Restore Points

Azure SQL Data Warehouse includes automated database snapshots that can be leveraged to recover or copy a data warehouse to a previous state. These snapshots support an eight-hour recovery point objective (RPO) and are available to be used for 7 days. If you require a faster RPO or you require your snapshot to be available for longer than 7 days, you can manually trigger a snapshot to save the current database state. This is good practice to follow before and after large modifications to your data

warehouse - it allows quicker recovery times in the event of any workload interruptions or user errors.

We will create a user-defined restore point, but not actually perform a restore due to the limitations of working in this lab environment and time constraints.

1. Logon to Azure Portal (portal.azure.com) using your credentials
2. Navigate to the Overview page of your Azure Data Warehouse.
3. Select New Restore Point



4. Specify a name for your restore point

User-Defined Restore Points □ ×

Create a new user-defined restore point.

User-defined restore points will have a default retention period of 7 days.

[Learn more](#) 

* Name

RestorePointOne 

OK

5. Click Ok
6. Navigate to the Overview page of your Azure Data Warehouse.
7. Select Restore



8. Select Restore Point Type: User-Defined Restore Points
Click Restore points to activate the User-Defined Restore Points blade

Notice your recently created user-defined restore point.

Dashboard > trichterusgsdatawarehouse (trichterusgsSqlServer/trichterusgsdatawarehouse) > Restore

Restore

There is a price implication when a new database is created.

Restore Point Type:
User-Defined Restore Points

* Database name
trichterusgsdatawarehouse_2019-02-05T20-20

User-Defined Restore Points

* Restore points
None

* Server
trichterusgsSqlServer (East US 2)

* Performance level ⓘ
Gen2: DW1000c

OK

User-Defined Restore Points

Create user-defined restore point

RestorePointOne
2019-02-05T20:15:50Z

9. Click X to close the User-Defined Restore Points blade

10. Investigate the Automatic Restore Points

Select Restore Point Type: Automatic Restore Points


Notice the Oldest restore point and Newest restore point information

Select the calendar control and click a **bold** date

Notice that the time drop down is now populated with the restore point times available.

Restore

□ ×



There is a price implication when a new database is created.

Restore Point Type:

Automatic Restore Points ▾

* Database name

trichterugsgdatawarehouse_2019-02-04T22- ✓


Automatic Restore Points

Current time
2019-02-05 21:17 UTC

Oldest restore point
2019-02-01 22:43 UTC

Newest restore point
2019-02-05 20:15 UTC

Current time
Restore points are created at least every 8 hours

2019-02-04 

22:43:41 ▾

UTC

Restore points are created at least every 8 hours

* Server

trichterugsgsqlserver (East US 2) >

* Performance level ⓘ

Gen2: DW1000c >

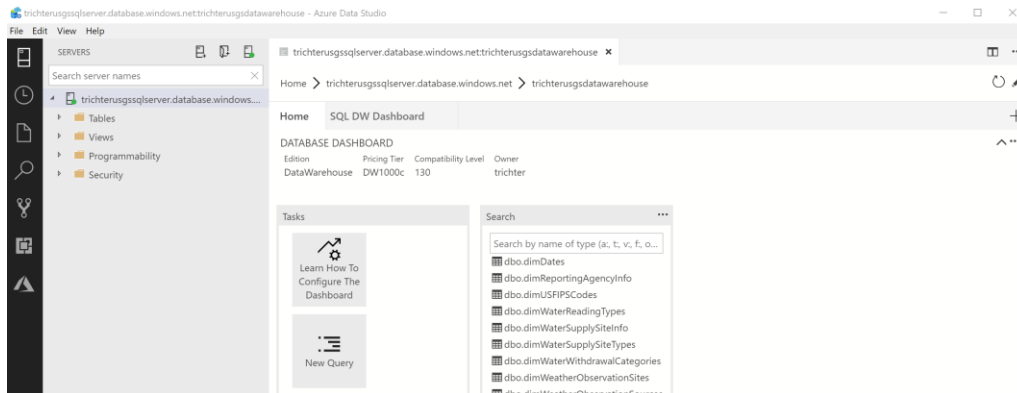
11. Click x to close the Restore blade.

We will not perform an actual restore due to time and resource constraints of this lab environment.

12. Let's see how restore points can be queried using Azure Data Studio

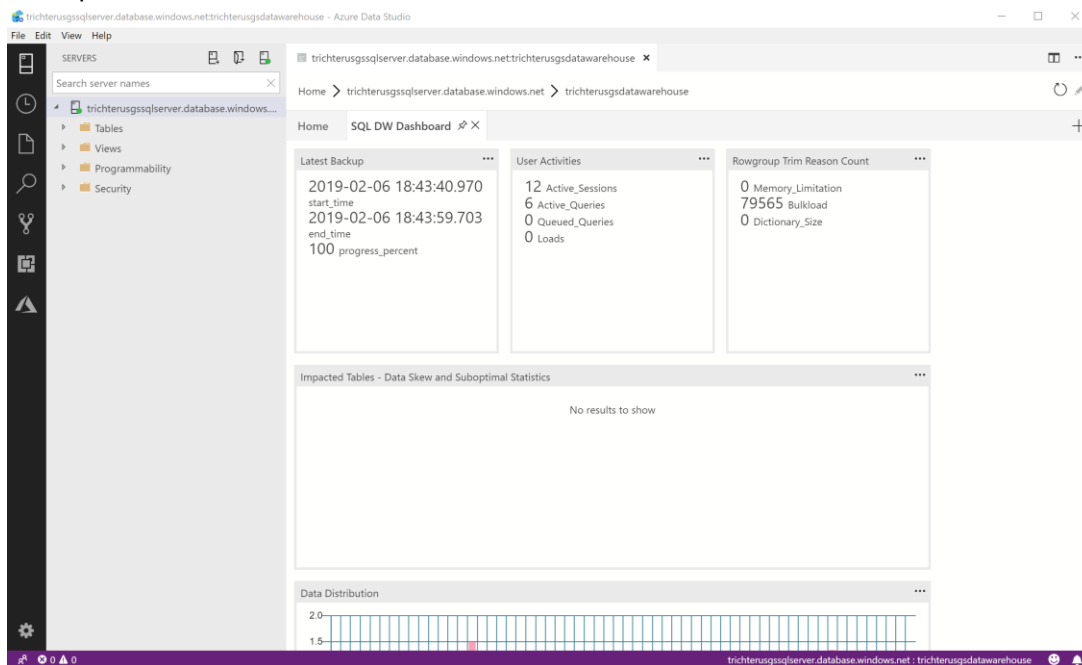
Open your Azure Data Studio application and connect to your Azure Data Warehouse
(You must have completed the prior module Azure SQL Data Studio SQLDW Dashboard)

13. Right click your Azure Data Warehouse server and select Manage




14. Select the SQL DW Dashboard tab

Notice the default summary information provided for user activities, CCI Health (Rowgroup Trim Reason Count), Tables that are being impacted by skew or by statistics issues, Etc. The queries that run behind these visualizations are available to you. Let's look into the Latest Backup information.

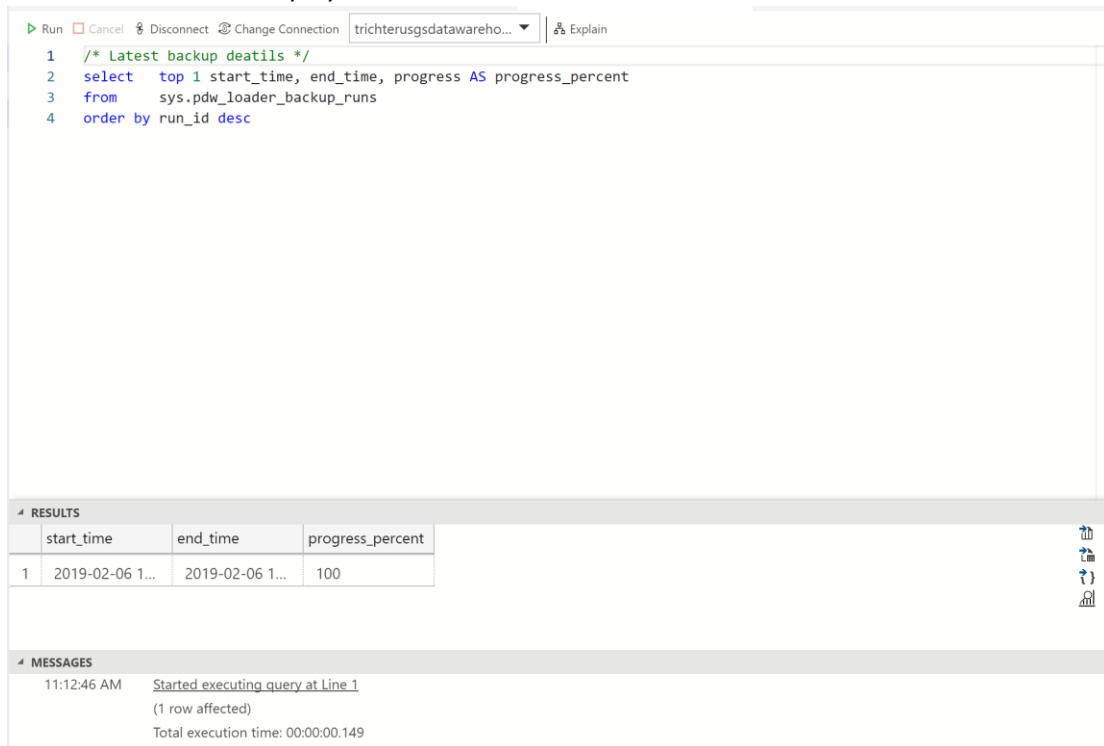


15. Select the ellipses (...) in the right-hand corner of the Latest Backup visualization



16. Select Run Query

A new tab will open and execute the query that populates the Latest Backup visualization and the result set will be displayed.



The screenshot shows a SQL query editor with a toolbar at the top containing 'Run', 'Cancel', 'Disconnect', 'Change Connection', and a dropdown menu set to 'trichterusgsdatawareho...'. The query text is as follows:

```
1  /* Latest backup deatils */
2  select top 1 start_time, end_time, progress AS progress_percent
3  from sys.pdw_loader_backup_runs
4  order by run_id desc
```

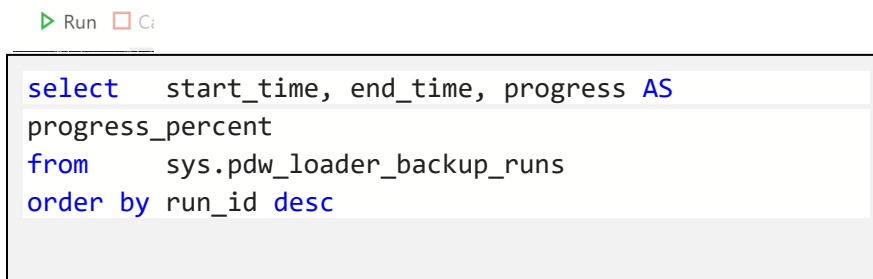
Below the query editor, the 'RESULTS' section displays a table with the following data:

	start_time	end_time	progress_percent
1	2019-02-06 1...	2019-02-06 1...	100

The 'MESSAGES' section at the bottom shows the following information:

11:12:46 AM Started executing query at Line 1
(1 row affected)
Total execution time: 00:00:00.149

17. To review all of the backups that are available, modify the query by removing the top 1 clause from the select statement and click run



The screenshot shows a SQL query editor with a toolbar at the top containing 'Run' and 'Cancel'. The query text is as follows:

```
select start_time, end_time, progress AS
progress_percent
from sys.pdw_loader_backup_runs
order by run_id desc
```

18. All of the available restore points will be listed. Notice that they do not have friendly names. Your user-defined restore point is likely at the top, check the date and time.

Run Cancel Disconnect Change Connection trichterugsdatabareho... Explain

```

1  /* Latest backup deatils */
2  select start_time, end_time, progress AS progress_percent
3  from sys.pdw_loader_backup_runs
4  order by run_id desc

```

RESULTS

	start_time	end_time	progress_percent
3	2019-02-05 20:15:50.610	2019-02-05 20:16:11.890	100
4	2019-02-04 22:43:41.010	2019-02-04 22:44:02.763	100
5	2019-02-04 18:43:41.013	2019-02-04 18:44:01.717	100
6	2019-02-04 14:43:41.010	2019-02-04 14:44:01.650	100
7	2019-02-04 10:43:41.020	2019-02-04 10:44:02.803	100
8	2019-02-04 06:43:40.997	2019-02-04 06:44:01.730	100
9	2019-02-04 02:43:41.027	2019-02-04 02:44:01.760	100
10	2019-02-03 22:43:41.020	2019-02-03 22:44:01.680	100
11	2019-02-03 18:43:41.030	2019-02-03 18:44:00.733	100
12	2019-02-03 14:43:41.007	2019-02-03 14:44:00.723	100
13	2019-02-03 10:43:41.017	2019-02-03 10:44:01.690	100

MESSAGES

11:12:46 AM Started executing query at Line 1
(1 row affected)
Total execution time: 00:00:00.149

11:13:41 AM Started executing query at Line 1
(22 rows affected)
Total execution time: 00:00:00.111

19. Close the tab. Don't Save your changes

Maintenance Window Scheduling, Service Health, Service Health Alerts

By default, all newly created Azure SQL Data Warehouse instances have an eight-hour primary and secondary maintenance window applied during deployment. You can change the windows as soon deployment is complete. No maintenance will take place outside the specified maintenance windows without prior notification.

1. Logon to Azure Portal (portal.azure.com) using your credentials
2. Navigate to your data warehouse overview blade

Notice the Maintenance schedule text link

The state of any maintenance will be displayed in this area. "Not yet active" indicates that the current day and time are not within an active maintenance window.

trichterugsdatabarehouse (trichterugssqlserver/trichterugsdatabarehouse)

SQL data warehouse

Search (Ctrl+F)

Pause Scale Restore New Restore Point Delete

Overview

Activity log

Tags

Diagnose and solve problems

Settings

Resource group (change) : trichterPreReadyLabs

Status : Online

Location : East US 2

Subscription (change) : SQL-D5-DW_SU06_R&D_60843

Subscription ID : f76b23dd-f29b-4562-a6c2-7299fe286f4b

Server name : trichterugssqlserver.database.windows.net

Connection strings : Show database connection strings

Performance level : Gen2: DW1000c

Maintenance schedule : Not yet active: Sat 06:00 UTC (5h) / Tue 22:00 UTC (8h)

Geo-backup policy : Enabled

3. Click on Not Yet Active text link to activate the Maintenance schedule blade

Dashboard > trichterugsdatawarehouse (trichterugssqlserver/trichterugsdatawarehouse) > Maintenance schedule

Save Discard Feedback

Information Maintenance schedule is not yet available in this region. However, the schedule you configure will automatically become active once the feature is available. →

Information Maintenance on your data warehouse could occur once a week within one of two maintenance windows. Choose the primary and secondary windows that best suit your operational needs. If you would like to use the maintenance windows already defined, no action is required. Maintenance will not take place outside these windows unless we notify you in advance.

Choose primary window

☒ Saturday - Sunday ☐ Tuesday - Thursday

Primary maintenance window	Secondary maintenance window
Day: <input type="text" value="Saturday"/>	Day: <input type="text" value="Tuesday"/>
Start time: <input type="text" value="06:00 UTC"/>	Start time: <input type="text" value="22:00 UTC"/>
Time window: <input type="text" value="5 hours"/>	Time window: <input type="text" value="8 hours"/>

Schedule summary

Primary maintenance window Saturday 06:00 UTC (5 hours)	Secondary maintenance window Tuesday 22:00 UTC (8 hours)
--	---

4. Notice the various configuration options and configure a Primary and Secondary maintenance window

Dashboard > trichterusgsdatawarehouse (trichterusgssqlserver/trichterusgsdatawarehouse) > Maintenance schedule

trichterusgsdatawarehouse

Save Discard Feedback

Maintenance schedule is not yet available in this region. However, the schedule you configure will automatically become active once the feature is available.

Maintenance on your data warehouse could occur once a week within one of two maintenance windows. Choose the primary and secondary windows that best suit your operational needs. If you would like to use the maintenance windows already defined, no action is required.

Maintenance will not take place outside these windows unless we notify you in advance.

Choose primary window

☐ Saturday - Sunday ☒ Tuesday - Thursday

Primary maintenance window	Secondary maintenance window
Day Wednesday	Day Saturday
Start time 05:00 UTC	Start time 06:00 UTC
Time window 3 hours	Time window 5 hours

Schedule summary


Primary maintenance window	Secondary maintenance window
Wednesday 05:00 UTC (3 hours)	Saturday 06:00 UTC (5 hours)

- Click Save
- Return to your Overview blade and notice the new days and times listed for Maintenance schedule


Performance level : Gen2: DW1000c

Maintenance schedule : **Not yet active: Wed 05:00 UTC (3h) / Sat 06:00 UTC (5h)**

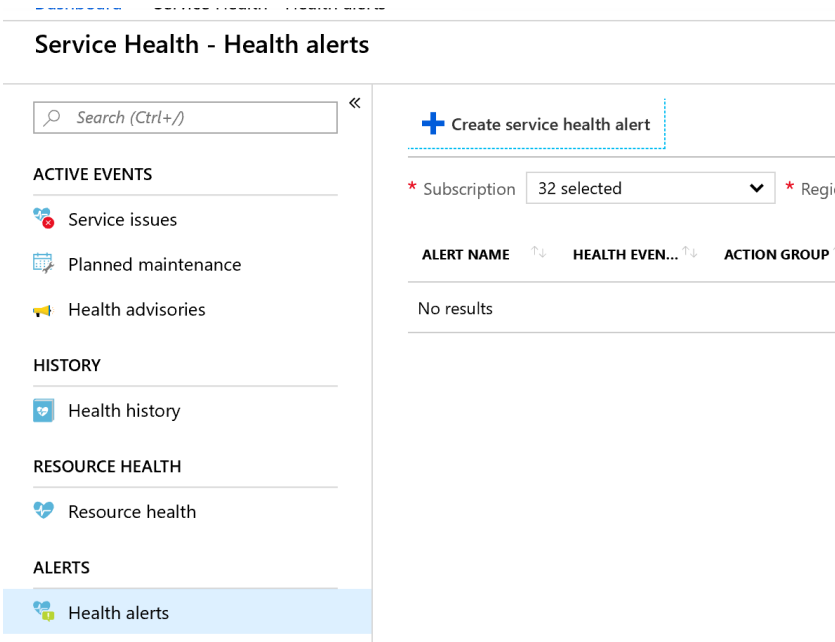
- Planned Maintenance Activities will appear in Service Health
From the Dashboard select Service Health



- Click Planned Maintenance from the side menu to view Planned Maintenance events and notifications
Filter for SQL Data Warehouse



9. To be alerted for Planned SQL Data Warehouse planned maintenance
Select Health alerts and click Create service health alert



10. Configure a new rule

Dashboard > Service Health - Health alerts > Create rule

Create rule

Rules management

Event type: Planned maintenance

ACTION GROUPS
 Notify your team via email and text messages or automate actions using webhooks, runbooks, functions, logic apps or integrating with external ITSM solutions. Learn more [here](#)

ACTION GROUP NAME	ACTION GROUP TYPE
trichterPreReadyADWmaintenance	1 Email

Select existing Create New

ALERT DETAILS

Alert rule name: trichterPreReadyLabsADWmaintenance ✓

Description: email me when the Pre-Ready lab ADW has a planned maintenance event ✓

Save alert to resource group: trichterPreReadyLabs ✓

Enable rule upon creation: Yes No

Create alert rule

11. Click Create alert rule and refresh the Health alerts blade

Dashboard > Service Health - Health alerts

Service Health - Health alerts

Search (Ctrl+F)

+ Create service health alert

Subscription: 32 selected Region: 31 selected Services: 146 selected Health Event Type: 3 selected

ALERT NAME	HEALTH EVENT	ACTION GROUP	RESOURCE G.	SUBSCRIPTION	LAST FIRED	STATUS
trichterPreReady...	Planned Maintenance	trichterpreready...	trichterPreReady...	7f6b236d-02b0-4562-af62-7299fe28...	Never in last two weeks	Enabled

Details History

Alert name
trichterPreReadyLabsADWmaintenance

Alert criteria
 Health event type: Planned Maintenance
 Subscription: 7f6b236d-02b0-4562-af62-7299fe2864b (SQL-DW-SUBS_RBD_60843)
 Region(s): eastus2
 Service(s): SQL Data Warehouse

Alert via
 Action group name: trichterPreReadyADWmaintenance [Edit this action group](#)

Was this helpful?

12. You will be able to view the details and history of this health alert from this blade.

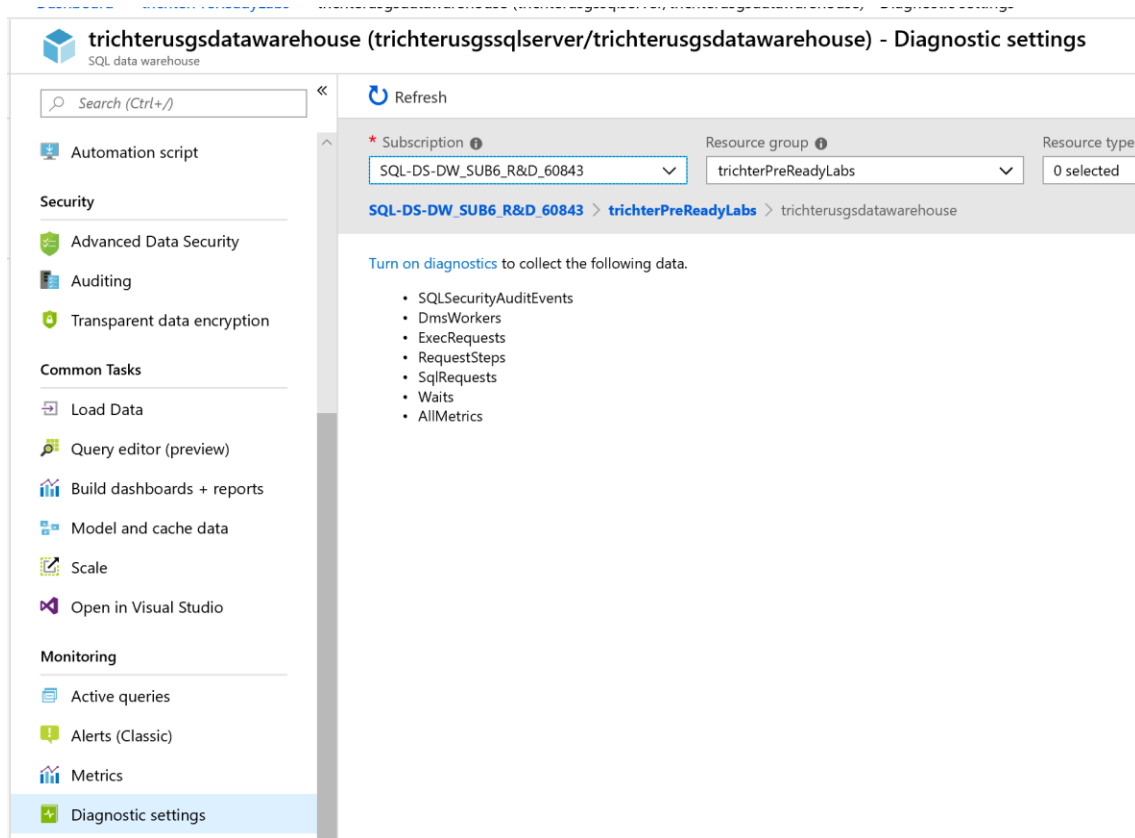
13. Disable your alert by clicking the ellipses(...) and selecting Disable

Querying ADW Diagnostic Logs using Azure Monitor

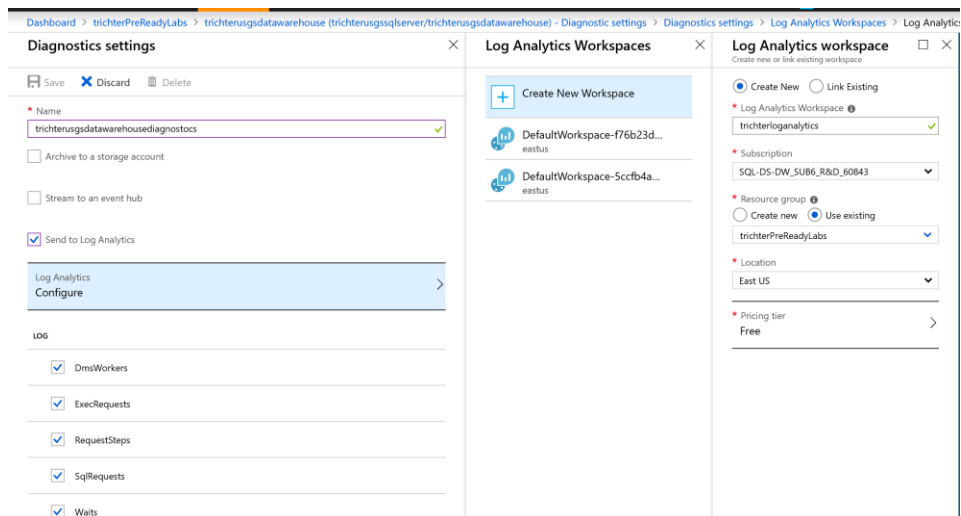
SQL Data Warehouse (SQL DW) now enables enhanced insights into analytical workloads by integrating directly with Microsoft Azure Monitor diagnostic logs. This new capability enables developers to analyze workload behavior over an extended time period and make informed decisions on query optimization or capacity management.

If you did not configure Diagnostics earlier:

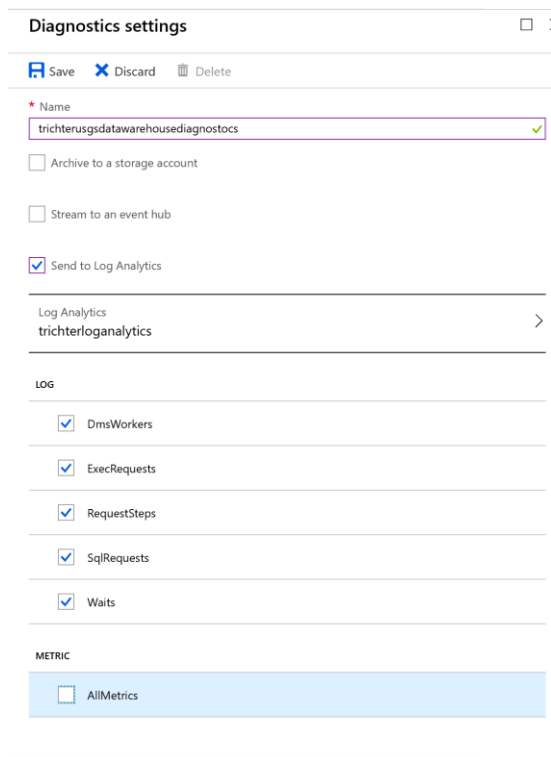
1. Logon to Azure Portal (portal.azure.com) using your credentials
2. Navigate to your Azure Data Warehouse
3. Click on Diagnostic Settings from the side menu



4. Click on the Turn on diagnostics text link in the blade
5. Provide a name for your diagnostics
 - Check Send to Log Analytics
 - Select all options in LOG
 - Click Log Analytics Configure to activate the Log Analytics Workspaces blade
 - Create a new Workspace in your resource group using one of the charged pricing tiers



6. Click Save on Diagnostics settings blade



7. Click refresh on Diagnostic settings blade to see your new diagnostics

Refresh

Subscription ⓘ

SQL-D5-DW_SUB6_R&D_60843

Resource group ⓘ

trichterPreReadyLabs

Resource type ⓘ

0 selected

Resource ⓘ

Type to start filtering ...

SQL-D5-DW_SUB6_R&D_60843

>

trichterPreReadyLabs

>

trichterusgsdatawarehouse

NAME	STORAGE ACCOUNT	EVENT HUB	LOG ANALYTIC	EDIT SETTING
trichterusgsdatawarehouseiagnostics	-	-	trichterloganalytics	<a>Edit setting

Add diagnostic setting

Review Diagnostics Logs:

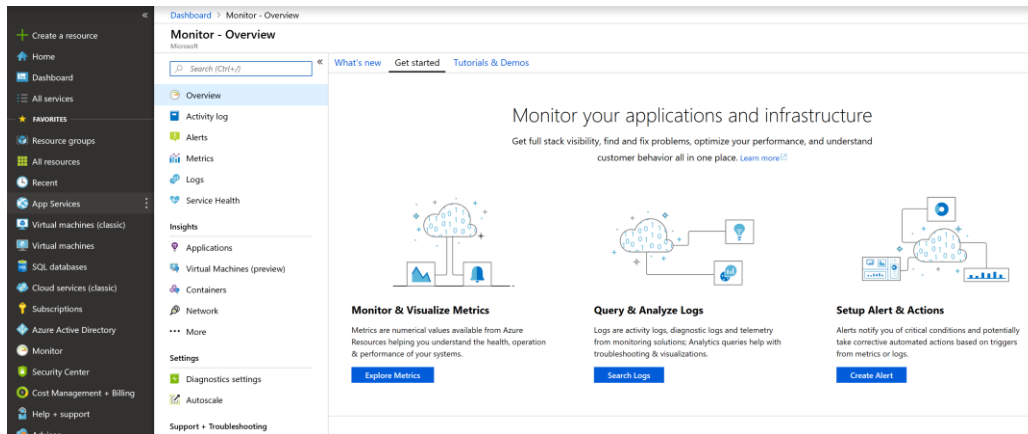
1. To collect some data, using SSMS connect to your Azure Data Warehouse and execute some queries

```
SELECT FIS.SalesAmount, DG.PostalCode, DC.YearlyIncome AS CustomerIncome,
DD.FullDateAlternateKey AS OrderDate
FROM FactInternetSales AS FIS
JOIN DimCustomer AS DC
ON (FIS.CustomerKey = DC.CustomerKey)
JOIN DimDate AS DD
ON (FIS.OrderDateKey = DD.DateKey)
```

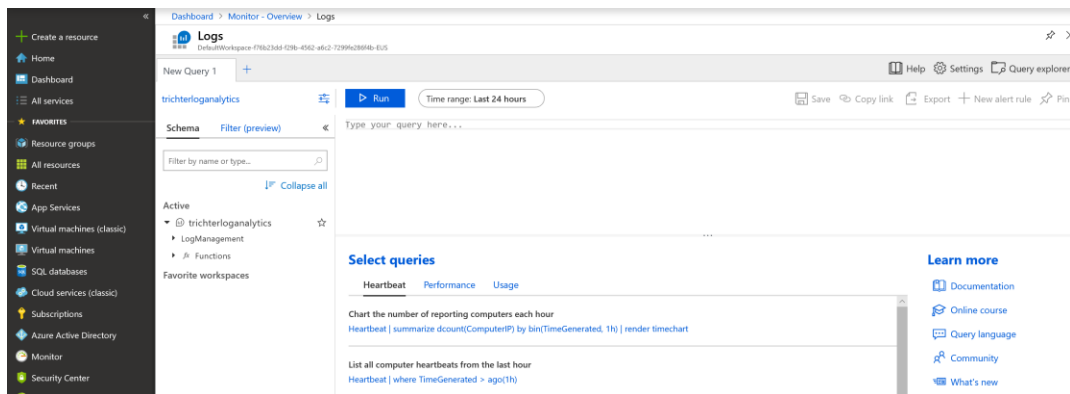
```
SELECT fis.SalesAmount, dst.Gender, dst.NumberCarsOwned, dst.YearlyIncome
AS CustomerYearlyIncome, dst.TotalChildren
FROM FactInternetSales AS fis
LEFT OUTER JOIN DimCustomer AS dst
ON (fis.CustomerKey=dst.CustomerKey);
```

```
SELECT fis.SalesAmount, dst.ProductLine
FROM FactInternetSales AS fis
LEFT OUTER JOIN DimProduct AS dst
ON (fis.ProductKey=dst.ProductKey);
```

2. In the portal select Monitor from the sidebar to activate the Monitor – Overview blade



3. Click Search Logs



4. Execute these sample queries to get a feel for the information that can be retrieved from Azure Diagnostics.

Chart to determine the most active resource classes by request

```
AzureDiagnostics
| where Category contains "ExecRequests"
| where Status_s == "Completed"
| summarize totalQueries = dcount(RequestId_s) by
ResourceClass_s
| render barchart
```

Count of all queued queries

```
AzureDiagnostics
| where Category contains "waits"
| where Type_s == "UserConcurrencyResourceType"
| summarize totalQueuedQueries =
dcount(RequestId_s)
```

Chart for top requests most impacted by data movement operations

```
AzureDiagnostics
| where Category == "RequestSteps"
| where OperationType_s in ("ShuffleMoveOperation",
"BroadcastMoveOperation", "PartitionMoveOperation",
"RoundRobinMoveOperation",
"SingleSourceRoundRobinMoveOperation",
"MoveOperation", "TrimMoveOperation")
| where Status_s == "Complete"
| project RequestId_s,
duration=datetime_diff('millisecond',EndTime_t,
StartTime_t)
| order by duration desc
| take 10
| render barchart
```